



Cooperative OTP Ring Simulation

Nithesh Gurudaas Kujuluva Ganesh, Raffael Davila, Andy Then

February 17, 2026

1 Overview

This document describes a simulation of a cooperative one-time pad (OTP) ring protocol operating over an asynchronous network. The simulation models m parties arranged in a logical ring, each maintaining a pointer into a shared pad of length n , subject to a minimum gap constraint d enforced between adjacent parties. The key question studied is: how many pads are wasted (i.e. left unused) when only x out of m parties are actively sending?

2 System Components

2.1 Asynchronous Network

The `AsynchronousNetwork` class simulates a broadcast medium with variable delivery delay. Each broadcast from a sender arrives at all other nodes within a uniformly random delay of $[0, d_{\text{delay}}]$ ticks. The global clock advances by one unit per call to `tick`, which delivers all messages whose scheduled arrival time has been reached and updates the recipient parties' views accordingly.

Formally, when party i broadcasts its new index idx at time t , each recipient $j \neq i$ receives the update at time:

$$t_{\text{deliver}} = t + \delta, \quad \delta \sim \mathcal{U}\{0, d_{\text{delay}}\}.$$

2.2 Ring Party

Each `RingParty` represents a node in the ring with the following state:

- `my_index`: the party's current position in the pad array, initialised to $(\text{party_id} - 1) \cdot \lfloor n/m \rfloor$.
- `view_of_others`: a dictionary mapping each party ID to its last known index, initialised to the same evenly-spaced starting positions.
- `pads_used`: a counter of pad indices this party has consumed for encryption.

The gap constraint requires that a party may only advance if the distance to its forward neighbor (modulo n) exceeds d :

$$\text{gap} = (\text{neighbor_pos} - \text{my_index}) \bmod n > d.$$

3 Simulation Protocol

3.1 Initialisation

Given parameters n , m , d , and x , the simulation proceeds as follows:

- x parties are designated *active* (chosen uniformly at random); the remaining $m - x$ are *silent*.
- A *burned* set tracks pad indices that have been consumed. It is initialised with the starting positions of all active parties.
- The theoretical maximum utilisation is $n - m \cdot d$, accounting for the mandatory gaps between all m parties.

3.2 Move Classification

At each tick, the eligibility of each party to advance is determined by `get_move_status`, which returns one of three outcomes:

`data` The gap constraint is satisfied and the next index is fresh (not burned). The party encrypts a message, burns the pad, and increments `pads_used`.

`drift` The gap constraint is satisfied but the next index is already burned. The party advances without consuming a new pad (skipping over a used index).

`None` The gap constraint is violated; the party is blocked and cannot move.

3.3 Scheduling Policy

Each tick applies a two-priority scheduling rule:

1. **Active parties have priority.** Any active party that is not blocked may move. One such party is chosen uniformly at random, advances its index, and broadcasts its new position.
2. **Silent parties move only when no active party can.** If all active parties are blocked, a random unblocked silent party performs a *yield* move: it advances without burning and broadcasts, thereby clearing space for active neighbours.

3.4 Termination

The simulation terminates when either:

- the burned set reaches the maximum utilisation $n - m \cdot d$, or
- no party can move *and* the network message queue is empty (*Clinch* state).

The function returns $n - |\text{burned}|$, the count of unused pad indices.

4 Algorithm

5 Experimental Setup

The simulation is run for $M \in \{3, 4\}$ parties with fixed parameters $N = 2000$ and $D = 15$. For each value of M , every scenario $S.x$ (for $x = 1, \dots, M$) is repeated over 50 independent trials. The reported metrics are:

- **Average wasted pads:** $\mathbb{E}[n - |\text{burned}|]$ across trials.
- **Utilisation:** $\left(1 - \frac{\mathbb{E}[\text{wasted}]}{n}\right) \times 100\%$.

Scenario $S.x$ denotes the case where exactly x out of M parties are active senders; the remaining $M - x$ parties are silent and only perform yield moves to vacate space.

Algorithm 1 Cooperative OTP Ring Simulation

1: **Input:** pad length n , party count m , gap bound d , active count x
2: **Output:** number of unused pad indices
3: Initialise NETWORK with delay d
4: Sample active set $A \subseteq \{1, \dots, m\}$, $|A| = x$; silent set $S = \{1, \dots, m\} \setminus A$
5: Initialise parties: $\text{pos}[i] \leftarrow (i - 1)\lfloor n/m \rfloor$ for all i
6: $\text{burned} \leftarrow \{\text{pos}[i] : i \in A\}$
7: $\text{MAX} \leftarrow n - m \cdot d$
8: **while** $|\text{burned}| < \text{MAX}$ **do**
9: NETWORK.TICK(parties) ▷ Deliver pending messages, update views
10: $\text{movers} \leftarrow \{i \in A : \text{GAPOK}(i)\}$
11: **if** $\text{movers} \neq \emptyset$ **then**
12: *i* $\leftarrow \text{RANDOMCHOICE}(\text{movers})$
13: *nxt* $\leftarrow (\text{pos}[i] + 1) \bmod n$
14: $\text{pos}[i] \leftarrow \text{nxt}$
15: **if** $\text{nxt} \notin \text{burned}$ **then**
16: $\text{burned} \leftarrow \text{burned} \cup \{\text{nxt}\}$ ▷ **data** move
17: **else**
18: ▷ **drift** move; pad already used
19: **end if**
20: NETWORK.BROADCAST(*i*, *nxt*)
21: **else**
22: $\text{jumpers} \leftarrow \{j \in S : \text{GAPOK}(j)\}$
23: **if** $\text{jumpers} \neq \emptyset$ **then**
24: *j* $\leftarrow \text{RANDOMCHOICE}(\text{jumpers})$
25: $\text{pos}[j] \leftarrow (\text{pos}[j] + 1) \bmod n$
26: NETWORK.BROADCAST(*j*, $\text{pos}[j]$) ▷ **yield** move
27: **else**
28: **if** NETWORK.QUEUEEMPTY() **then**
29: **break** ▷ Clinch state reached
30: **end if**
31: **end if**
32: **end if**
33: **end while**
34: **return** $n - |\text{burned}|$
35: **procedure** $\text{GAPOK}(i)$
36: $\text{fwd} \leftarrow (i \bmod m) + 1$
37: **return** $(\text{view}[i][\text{fwd}] - \text{pos}[i]) \bmod n > d$
38: **end procedure**

6 Security Property

The simulation includes a *loud-fail* check: if an active party attempts a **data** move onto an index already present in **burned**, a RuntimeError is raised immediately. This guarantees that pad reuse is detectable at simulation time.